



**Tanta University**

**Faculty of Computers and Information**

**Information Technology Department**

## **LSTM Report for Multimedia Mining Oral Test**

**Supervised by\ Dr. Elhossiny Ibrahim**

**Prepared by: Ibrahim Mohamed Ibrahim Eita. -> Section: 9.**

LSTM stands for Long Short-Term Memory, which is a type of recurrent neural network (RNN) architecture. RNNs, including LSTMs, are designed to process sequential data, such as time series data or natural language sentences, where the order of the elements matters.

LSTMs were introduced to address the limitations of traditional RNNs in capturing long-term dependencies and dealing with the vanishing gradient problem.

The vanishing gradient problem refers to the issue where the gradients used for updating the weights in the network become very small, causing the network to have difficulty learning from and propagating information over long sequences.

LSTMs overcome the vanishing gradient problem by introducing a memory cell and different types of gates. The memory cell is responsible for storing and updating the information over time. The gates, including the input gate, forget gate, and output gate, control the flow of information into and out of the memory cell.

The input gate determines which parts of the input sequence are relevant to update the memory cell. The forget gate decides which information to discard from the memory cell. The output gate determines which parts of the memory cell content should be outputted to the next step in the sequence.

By using these gates and memory cell, LSTMs can selectively remember or forget information over long sequences, allowing them to capture and retain important information over time. This makes LSTMs particularly useful for tasks such as speech

recognition, language modeling, machine translation, sentiment analysis, and many other applications involving sequential data.

LSTMs are designed to address the limitations of traditional RNNs, such as the vanishing gradient problem, and they excel at capturing long-term dependencies in sequential data.

### **The key components of LSTM are as follows:**

- **Memory Cell:** This is the core component of LSTM and stores information over time. It allows LSTMs to maintain a memory state and control the flow of information.
- **Input Gate:** The input gate regulates the flow of new information into the memory cell. It decides which parts of the input sequence are relevant and should be stored.
- **Forget Gate:** The forget gate determines which information to discard from the memory cell. It controls the extent to which previous information is retained or forgotten.
- **Output Gate:** The output gate selects which parts of the memory cell content should be outputted to the next step in the sequence. It influences the output of the LSTM at each time step.

LSTMs are trained using the backpropagation through time (BPTT) algorithm, which extends backpropagation for recurrent networks.

BPTT calculates gradients by unfolding the network over time and propagating errors through each time step, enabling the LSTM to learn and adjust its parameters based on the temporal relationships in the training data.

### **There are also variations of LSTM that have been introduced over time:**

- **Gated Recurrent Unit (GRU):** GRU is a simplified version of LSTM that combines the forget and input gates into a single update gate. It aims to solve similar problems as LSTM but with a slightly different architecture.
- **Bidirectional LSTM (BiLSTM):** BiLSTMs process sequences in both forward and backward directions, capturing information from past and future contexts. They are useful for tasks that require a comprehensive understanding of the input sequence.
- **Peephole Connections:** Peephole connections extend LSTMs by allowing the gates to have direct access to the memory cell. This enables the gates to consider the memory content when making decisions.

LSTMs have found applications in various domains, including Natural Language Processing (NLP), speech recognition, time series analysis, and image/video analysis.

They are used for tasks such as language modeling, sentiment analysis, machine translation, stock price forecasting, video classification, and more.

Several popular deep learning frameworks, such as TensorFlow, PyTorch, Keras, and Caffe, provide implementations of LSTM and related architectures, making it easier to work with LSTMs and incorporate them into different projects.

### **Handling Long-Term Dependencies:**

LSTMs excel at capturing long-term dependencies in sequential data.

Unlike traditional RNNs, which struggle with the vanishing gradient problem, LSTMs effectively propagate and retain information over long sequences.

This is accomplished by the memory cell and the gating mechanisms, which selectively remember and forget information.

### **Training LSTMs:**

LSTMs are typically trained using the backpropagation through time (BPTT) algorithm, an extension of backpropagation for recurrent networks.

BPTT calculates gradients by unfolding the network over time and propagating errors through each time step.

This allows the LSTM to learn and adjust its parameters based on the temporal relationships in the training data.

### **LSTMs have found success in various domains and applications, including:**

- **Natural Language Processing (NLP):** LSTMs are used for tasks such as language modeling, sentiment analysis, machine translation, text generation, and named entity recognition.
- **Speech Recognition:** LSTMs can model the temporal dependencies in audio data and are employed in speech recognition systems.
- **Time Series Analysis:** LSTMs are effective for analyzing and forecasting time series data, such as stock prices, weather patterns, and energy consumption.
- **Image and Video Analysis:** LSTMs can be applied to analyze sequences of images or video frames, allowing tasks like video classification, action recognition, and video captioning.

- **Tools and Frameworks:** There are several popular deep learning frameworks that provide implementations of LSTM and related architectures, including TensorFlow, PyTorch, Keras, and Caffe.

## LSTM Model Building:

```
import numpy as np
from keras.models import Sequential
from keras.layers import LSTM, Dense, Embedding
from keras.preprocessing import sequence
from keras.datasets import imdb

# Set the parameters
max_features = 25000 # Number of words to consider as features
max_len = 100 # Maximum length of review (in words)
batch_size = 32

# Load the dataset
(X_train, y_train), (X_test, y_test) = imdb.load_data(num_words=max_features)
# Pad sequences to a fixed length
X_train = sequence.pad_sequences(X_train, maxlen=max_len)
X_test = sequence.pad_sequences(X_test, maxlen=max_len)
# Build the LSTM model
model = Sequential()
model.add(Embedding(max_features, 32))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))
# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# Train the model
model.fit(X_train, y_train, batch_size=batch_size, epochs=5, validation_data=(X_test,
y_test))
# Evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))
# Save the model
model.save("sentiment_analysis_model.h5")
```

This code snippet showcases the implementation of a sentiment analysis model using a recurrent neural network (RNN) architecture known as Long Short-Term Memory (LSTM). The code begins by importing the necessary libraries, including NumPy for numerical operations, as well as various components from the Keras deep learning framework.

Next, several parameters are set to configure the model and data preprocessing. The `max_features` parameter determines the maximum number of words to consider as features in the dataset, while `max_len` specifies the maximum length of a movie review in

terms of words. Additionally, `batch_size` determines the number of samples that will be propagated through the network during each training iteration.

The IMDB movie review dataset is then loaded using the `imdb.load_data()` function. This dataset is divided into training and testing sets, with the reviews and corresponding sentiment labels stored in `X_train`, `y_train`, `X_test`, and `y_test`, respectively. The `num_words` argument is set to `max_features`, which limits the vocabulary size to the most frequently occurring `max_features` words in the dataset.

To ensure uniformity in the length of reviews, the sequences are preprocessed using `sequence.pad_sequences()`. Both the training and testing reviews are padded or truncated to a fixed length of `max_len` words. This step is crucial for maintaining consistent input sizes for the LSTM model.

The LSTM model is then constructed using the `Sequential` class from Keras. The model architecture consists of an embedding layer, an LSTM layer, and a dense layer. The embedding layer maps the words in the reviews to dense vectors of a fixed size (32 in this case). The LSTM layer, with 32 units, captures temporal dependencies in the review sequences. Lastly, a dense layer with a single unit and sigmoid activation function is added to obtain a binary sentiment prediction, determining whether the review is positive or negative.

Once the model architecture is defined, the code proceeds to compile the model. The compilation involves specifying the loss function, optimizer, and evaluation metric. In this case, the binary cross-entropy loss function (`binary_crossentropy`), Adam optimizer, and accuracy metric are utilized.

The model is then trained using the `fit()` function, which takes the training data, batch size, number of epochs (set to 5 in this instance), and the testing data for validation during the training process. This step allows the model to learn from the training data and optimize its parameters to make accurate sentiment predictions.

After training, the model's performance is evaluated on the testing set using the `evaluate()` function. The evaluation results, including the accuracy, are obtained and printed to assess the model's effectiveness in sentiment analysis.

Finally, the trained model is saved to a file named "`sentiment_analysis_model.h5`" using the `save()` method. This enables the model to be reused or deployed for sentiment analysis tasks in the future.

## Model Deployment:

```
from keras.models import load_model
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences

# Load the saved model
model = load_model("sentiment_analysis_model.h5")

# Load the file containing statements
file_path = "reviews.txt"

# Preprocess the statements
max_features = 5000 # Number of words to consider as features
max_len = 100 # Maximum length of sequences

# Load the statements from the file
with open(file_path, "r") as file:
    statements = file.readlines()

# Initialize tokenizer
tokenizer = Tokenizer(num_words=max_features)
tokenizer.fit_on_texts(statements)

# Convert text to sequences
sequences = tokenizer.texts_to_sequences(statements)

# Pad sequences to a fixed length
padded_sequences = pad_sequences(sequences, maxlen=max_len)

# Make predictions
predictions = model.predict(padded_sequences)

# Interpret predictions
sentiments = ["positive" if prediction > 0.5 else "negative" for prediction in
predictions]

# Write the output to a file
predictions = "predictions.txt"
with open(predictions, "w") as output_file:
    for statement, sentiment in zip(statements, sentiments):
        output_file.write("Statement: {}\n".format(statement.strip()))
        output_file.write("Predicted Sentiment: {}\n".format(sentiment))
        output_file.write("-----\n")
```

This code snippet demonstrates a detailed procedure for utilizing a pre-trained sentiment analysis model to predict sentiments of statements stored in a file. The code starts by loading the saved model using the `load_model()` function from the Keras library, enabling the retrieval of the trained model's architecture and weights. This allows for the utilization of the model's sentiment analysis capabilities without the need for retraining.

Next, the code specifies the file path and preprocessing parameters required for handling the statements. The `max_features` parameter determines the maximum number of words to consider as features, while `max_len` sets the maximum length of the sequences. These parameters play a crucial role in ensuring consistency during tokenization and sequence padding.

The statements are read from the file specified by `file_path` using the `open()` function in read mode. The `readlines()` method is employed to read all the lines in the file and store them as individual statements in the `statements` list. This step is crucial for accessing and processing each statement individually.

To prepare the text data for analysis, the code initializes a tokenizer object using the `Tokenizer` class from Keras. The `num_words` parameter is set to `max_features`, which defines the maximum number of words to consider in the tokenization process. The tokenizer is then fitted on the `statements` list using the `fit_on_texts()` method. This step ensures that the tokenizer learns the vocabulary based on the provided statements, enabling the conversion of text to sequences.

The code proceeds by converting the statements into sequences using the `texts_to_sequences()` method of the tokenizer. This conversion replaces each word in the statements with its corresponding index in the tokenizer's word index. The resulting sequences are stored in the `sequences` list, preserving the sequential structure of the statements.

To ensure uniformity in sequence length, the sequences are padded or truncated to a fixed length of `max_len` using the `pad_sequences()` function. Padding involves adding zeros to the beginning or end of shorter sequences, while truncation removes excess words from longer sequences. This step is essential for maintaining consistent input dimensions when making predictions using the pre-trained model.

Once the sequences are appropriately processed, the pre-trained model is utilized to predict the sentiments of the statements. The `predict()` method is applied to the padded sequences, generating a list of prediction values. These values represent the probability of a statement being positive or negative.

To interpret the predictions, the code iterates over each prediction value in the list and assigns a sentiment label accordingly. If the prediction value exceeds 0.5, the sentiment is considered positive; otherwise, it is considered negative. The interpretations are stored in the sentiments list, reflecting the sentiment prediction for each corresponding statement.

Finally, the code writes the output to an output file named "`predictions.txt`". The file is opened in write mode, and a for loop is employed to iterate over the statements and sentiments simultaneously using the `zip()` function. Within each iteration, the statement and its predicted sentiment are written to the output file using the `write()` method. Additionally, a separator is included to improve the readability and organization of the output.

## **Sample Reviews & Output Predictions:**

**Statement:** I absolutely loved the movie! It was captivating from start to finish.

➔ **Predicted Sentiment: positive**

**Statement:** The acting in the film was outstanding. Each character delivered a remarkable performance.

➔ **Predicted Sentiment: positive**

**Statement:** The cinematography was lackluster and uninspiring. The visuals did not add anything to the overall experience.

➔ **Predicted Sentiment: positive**

**Statement:** This is the best romantic comedy I have ever seen. It made me laugh and cry.

➔ **Predicted Sentiment: positive**

**Statement:** The soundtrack of the film was awful. It did not match the mood of the scenes at all.

➔ **Predicted Sentiment: positive**

**Statement:** The cinematography in this movie is breathtaking. Every scene is visually stunning.

➔ **Predicted Sentiment: negative**

**Statement:** The storyline was engaging, and the plot twists kept me on the edge of my seat.



➔ **Predicted Sentiment: positive**

**Statement:** The soundtrack of the film was amazing. It perfectly complemented the mood of each scene.

➔ **Predicted Sentiment: positive**

**Statement:** The film had a confusing and muddled message. It left me feeling unsatisfied and disappointed.

➔ **Predicted Sentiment: negative**

**Statement:** The chemistry between the lead actors was electric. Their performances were incredibly believable.

➔ **Predicted Sentiment: positive**

**Statement:** This movie was a disappointment. It failed to live up to the hype and left me unsatisfied.

➔ **Predicted Sentiment: positive**

**Statement:** I was completely mesmerized by the visual effects. They were top-notch and added an extra layer of excitement.

➔ **Predicted Sentiment: negative**

**Statement:** The story was predictable and unoriginal. I could guess the ending within the first few minutes.

➔ **Predicted Sentiment: negative**